

# DiffKV Active Runtime — System Architecture

Serving stack drives the MLX inference core; the patched attention layer routes prefill vs. decode

