

DiffKV Active Runtime — System Architecture

Serving stack (top) drives the MLX inference core (bottom); the patched attention routes prefill vs. decode

