


Prefill (chunk = 512 tokens, repeated)



prefill → decode boundary: `mx.eval` · `mx.clear_cache` · resolve decode policy
compressed decode → drop the native prefill cache (footprint = DiffKV store only)

Decode (per token, per layer)

