


# Cache Lifecycle & Execution Flow

Chunked prefill (exact) → streaming compression → memory release → fused sparse decode

## PREFILL (chunk = 512 tokens, repeated)



**PREFILL → DECODE boundary: `mx.eval · mx.clear_cache · resolve` adaptive policy  
( $N \geq 16k \rightarrow$  drop native KVCache; decode footprint = DiffKV store only)**

## DECODE (per token, per layer)

