


Cache Lifecycle & Execution Flow

Chunked prefill (exact) → streaming compression → memory release → fused sparse decode

PREFILL (chunk = 512 tokens, repeated)



PREFILL → DECODE boundary: `mx.eval · mx.clear_cache · drop native prefill KV`
→ release peak GQA activations; decode footprint = DiffKV store only

DECODE (per token, per layer)

