


Prefill (chunk = 512 tokens, repeated)



**prefill → decode boundary: mx.eval · mx.clear_cache · resolve decode policy
compressed decode → drop the native prefill cache (footprint = DiffKV store only)**

Decode (per token, per layer)

