


Cache Lifecycle & Execution Flow

Chunked prefill (sparse) → streaming compression → memory release → fused sparse decode

PREFILL (chunk = 512 tokens, repeated)



PREFILL → DECODE boundary: `mx.eval` · `mx.clear_cache` · resolve decode policy
(compressed → drop native KVCache; decode footprint = DiffKV store only)

DECODE (per token, per layer)

