

native_triton_sparse_attn_decode(q , blocks, routing)

HAS_TRITON ?

no

Fallback path

_pytorch_vectorized_sparse_attn_decode()

DKV_TRITON_STRICT = 1 ?

yes

hard error
(re-raise)

no

PyTorch output
+ telemetry log

*strict mode is recommended for GPU regression testing:
a wrong-but-non-throwing kernel is otherwise invisible*

*applied in place — before the block's
softmax denominator is finalized, so
no ghost entry is added. K and V use
separate position arrays: the highest-
error rows differ between K and V .*

same output interface

@triton.autotune over (S_MAX, BLOCKS_PER_CHUNK, num_warps)
first five real-shape calls benchmarked → fastest config locked in

_fused_sparse_decode_kernel — one launch

① anchor score $s_{anc} = (q \cdot a_k) \cdot \text{scale}$

② delta projection $\tilde{q} = q V_K^\top \in \mathbb{R}^{32}$

③ token scores $\delta s_i = \tilde{q} \cdot U_i$ (all i in the block)

④ scatter $q \cdot R_K$ into the score at res_K_positions

⑤ scatter $p \cdot R_V$ into the numerator at res_V_positions

⑥ online softmax → accumulate O , lse

$O(KrB)$ — the block is never expanded to $\hat{K}$

flash-merge (log-sum-exp) → o